

Infrastructure as Code

Containers, Docker, and Docker Compose

Jason Kuruzovich

2026-01-16

Infrastructure as Code

System Definition and Reproducibility

Today's Agenda

Learning Objectives

1. Understand the **problems** that containers solve
2. Master **Docker fundamentals** - images, containers, volumes
3. Write effective **Dockerfiles** for Node.js applications
4. Use **Docker Compose** for multi-service orchestration
5. Implement **development workflows** with containers
6. Understand container **networking** and **data persistence**

The “Works on My Machine” Problem

Without Containers

Developer A: Node 18, MongoDB 6.0

Developer B: Node 20, MongoDB 7.0

Production: Node 16, MongoDB 5.0

"But it works on my machine!"

- Different OS versions
- Different dependencies
- Different configurations

- Different data
- Hours debugging environment issues

With Containers

All environments run identical containers

Developer A: Docker

Developer B: Docker

Production: Docker

"It works everywhere!"

- Identical runtime
- Identical dependencies
- Identical configuration
- Reproducible builds
- Focus on building features

Container Fundamentals

What is a Container?

A container is a **lightweight, standalone, executable package** that includes everything needed to run software:

- Code
- Runtime
- System tools
- Libraries
- Settings

...


Container

App Code	Runtime (Node.js)	Libs (npm pkgs)	Config (.env)
-------------	----------------------	--------------------	------------------

Isolated process(es)

Host Operating System


Containers vs. Virtual Machines



Key Differences

Aspect	Virtual Machines	Containers
Size	GBs	MBs
Startup	Minutes	Seconds
Isolation	Complete OS	Process-level
Performance	~5-10% overhead	Near-native
Density	10s per host	100s per host
Use Case	Different OS needs	Application packaging

Docker Architecture



Core Concepts

Image

- Read-only template
- Built from Dockerfile
- Composed of layers
- Versioned with tags
- Stored in registries

```
docker pull node:20-alpine
docker images
```

Container

- Running instance of image
- Isolated environment
- Has its own filesystem
- Can be started/stopped
- Ephemeral by default

```
docker run -it node:20-alpine
docker ps
```

Image Layers

Layer 5: COPY . .	Your code
Layer 4: RUN npm install	Dependencies
Layer 3: COPY package*.json ./	Package files
Layer 2: WORKDIR /app	Directory
Layer 1: FROM node:20-alpine	Base image

Each layer is cached - rebuild only changed layers

. . .

Key Insight: Order your Dockerfile commands from least to most frequently changing for optimal caching.

Writing Dockerfiles

Dockerfile Basics

```
# Start from a base image
FROM node:20-alpine

# Set working directory
WORKDIR /app

# Copy package files first (for caching)
COPY package*.json ./

# Install dependencies
RUN npm ci --only=production

# Copy application code
COPY . .

# Expose the port
EXPOSE 3000

# Define the command to run
CMD ["node", "server.js"]
```

Dockerfile Instructions

Instruction	Purpose	Example
FROM	Base image	FROM node:20-alpine
WORKDIR	Set working directory	WORKDIR /app
COPY	Copy files from host	COPY . .
ADD	Copy + extract archives	ADD app.tar.gz /app
RUN	Execute command (build time)	RUN npm install
CMD	Default command (run time)	CMD ["npm", "start"]

Instruction	Purpose	Example
ENTRYPOINT	Container entry point	ENTRYPOINT ["node"]
ENV	Environment variable	ENV NODE_ENV=production
EXPOSE	Document port	EXPOSE 3000
VOLUME	Mount point	VOLUME ["/data"]

Development vs. Production Dockerfile

Development

```
FROM node:20-alpine

WORKDIR /app

# Install all dependencies
COPY package*.json .
RUN npm install

# Copy source (or use volume)
COPY . .

# Development command with hot reload
CMD ["npm", "run", "dev"]
```

Production

```
FROM node:20-alpine AS builder
WORKDIR /app
COPY package*.json .
RUN npm ci
COPY . .
RUN npm run build

FROM node:20-alpine
WORKDIR /app
COPY --from=builder /app/dist ./dist
COPY --from=builder /app/node_modules ./node_modules
```

```
USER node
CMD ["node", "dist/server.js"]
```

Multi-Stage Builds

```
# Stage 1: Build
FROM node:20-alpine AS builder

WORKDIR /app
COPY package*.json ./
RUN npm ci
COPY . .
RUN npm run build

# Stage 2: Production
FROM node:20-alpine AS production

WORKDIR /app

# Copy only what we need from builder
COPY --from=builder /app/dist ./dist
COPY --from=builder /app/package*.json ./
RUN npm ci --only=production

# Security: non-root user
USER node

EXPOSE 3000
CMD ["node", "dist/server.js"]
```

Benefits: Smaller images, no build tools in production

Best Practices for Dockerfiles

1. Use specific tags - node:20-alpine not node:latest
2. Minimize layers - Combine related RUN commands
3. Order for caching - Less changing → more changing
4. Use .dockerignore - Exclude unnecessary files
5. Don't run as root - Add USER node

6. **Multi-stage builds** - Separate build and runtime
7. **Health checks** - Add `HEALTHCHECK` instruction

. . .

```
# .dockerignore
node_modules
npm-debug.log
.git
.env
*.md
```

Docker Commands

Essential Commands

```
# Build an image
docker build -t myapp:1.0 .

# List images
docker images

# Run a container
docker run -d -p 3000:3000 --name myapp myapp:1.0

# List running containers
docker ps

# List all containers
docker ps -a

# View logs
docker logs myapp
docker logs -f myapp # Follow

# Stop container
docker stop myapp

# Remove container
docker rm myapp
```

```
# Remove image
docker rmi myapp:1.0
```

Interactive Commands

```
# Run interactive shell
docker run -it node:20-alpine /bin/sh

# Execute command in running container
docker exec -it myapp /bin/sh

# Execute specific command
docker exec myapp npm run migrate

# Copy files to/from container
docker cp myapp:/app/logs ./logs
docker cp ./config.json myapp:/app/
```

Port Mapping

Host Machine	Container
--------------	-----------

:8080	:3000
-p 8080:3000	

```
docker run -p 8080:3000 myapp
```

```
# Multiple ports
```

```
docker run -p 8080:3000 -p 9229:9229 myapp
```

```
# All interfaces
```

```
docker run -p 0.0.0.0:8080:3000 myapp
```

```
# Random host port
```


```
docker run -p 3000 myapp
```

Volumes: Data Persistence

```
# Named volume (Docker-managed)
docker run -v mydata:/app/data myapp

# Bind mount (host directory)
docker run -v $(pwd)/src:/app/src myapp

# Read-only mount
docker run -v $(pwd)/config:/app/config:ro myapp
```



Docker Compose

What is Docker Compose?

Docker Compose is a tool for defining and running **multi-container applications**.

```
# docker-compose.yml
services:
  web:
    build: ./web
    ports:
      - "3000:3000"

  api:
```

```

build: ./api
ports:
- "4000:4000"

database:
image: mongo:7
volumes:
- mongo-data:/data/db

volumes:
mongo-data:

```

One command to start everything: `docker compose up`

Compose File Structure

```

# docker-compose.yml
version: "3.9" # Optional in recent versions

services:          # Container definitions
  service_name:
    image: ...
    build: ...
    ports: ...
    volumes: ...
    environment: ...
    depends_on: ...

volumes:          # Named volumes
  volume_name:

networks:         # Custom networks
  network_name:

secrets:          # Sensitive data
  secret_name:

```

A Complete MERN Compose File

```
services:
  # React Frontend
  frontend:
    build:
      context: ./frontend
      dockerfile: Dockerfile.dev
    ports:
      - "3000:3000"
    volumes:
      - ./frontend/src:/app/src
    environment:
      - REACT_APP_API_URL=http://localhost:4000
  depends_on:
    - api

  # Express API
  api:
    build:
      context: ./api
      dockerfile: Dockerfile.dev
    ports:
      - "4000:4000"
    volumes:
      - ./api/src:/app/src
    environment:
      - MONGODB_URI=mongodb://mongo:27017/myapp
      - NODE_ENV=development
  depends_on:
    - mongo

  # MongoDB Database
  mongo:
    image: mongo:7
    ports:
      - "27017:27017"
    volumes:
      - mongo-data:/data/db
    environment:
      - MONGO_INITDB_DATABASE=myapp
```

```
volumes:  
  mongo-data:
```

Service Configuration Options

```
services:  
  api:  
    # Build configuration  
    build:  
      context: ./api  
      dockerfile: Dockerfile  
      args:  
        NODE_VERSION: 20  
  
    # Or use existing image  
    image: node:20-alpine  
  
    # Container name  
    container_name: myapp-api  
  
    # Port mapping  
    ports:  
      - "4000:4000"  
      - "9229:9229" # Debug port  
  
    # Volume mounts  
    volumes:  
      - ./api:/app  
      - /app/node_modules # Anonymous volume  
  
    # Environment variables  
    environment:  
      - NODE_ENV=development  
      - DB_HOST=mongo  
  
    # Or from file  
    env_file:  
      - .env.development  
  
  # Dependencies
```

```
depends_on:
  - mongo
  - redis

# Restart policy
restart: unless-stopped

# Resource limits
deploy:
  resources:
    limits:
      cpus: '0.5'
      memory: 512M
```

Docker Compose Commands

```
# Start all services
docker compose up

# Start in detached mode (background)
docker compose up -d

# Start specific service
docker compose up api

# Stop all services
docker compose down

# Stop and remove volumes
docker compose down -v

# Rebuild images
docker compose build
docker compose up --build

# View logs
docker compose logs
docker compose logs -f api

# Execute command in service
```

```
docker compose exec api npm run migrate

# List running services
docker compose ps

# Scale a service
docker compose up -d --scale api=3
```

Networking in Compose

```
services:
  frontend:
    networks:
      - frontend-net

  api:
    networks:
      - frontend-net
      - backend-net

  database:
    networks:
      - backend-net

networks:
  frontend-net:
  backend-net:
```

frontend-net

frontend api

backend-net

database

Services can reach each other by name within a network.

Health Checks

```
services:
  api:
    build: ./api
    healthcheck:
      test: ["CMD", "curl", "-f", "http://localhost:4000/health"]
      interval: 30s
      timeout: 10s
      retries: 3
      start_period: 40s

  mongo:
    image: mongo:7
    healthcheck:
      test: ["CMD", "mongosh", "--eval", "db.adminCommand('ping')"]
      interval: 10s
      timeout: 5s
      retries: 5

  api-dependent:
    build: ./dependent
    depends_on:
      api:
        condition: service_healthy
```

Development Workflow

Hot Reloading with Volumes

```
services:
  api:
    build:
      context: ./api
```

```
    dockerfile: Dockerfile.dev
    volumes:
        # Mount source code
        - ./api/src:/app/src
        # Preserve node_modules from container
        - /app/node_modules
    command: npm run dev # nodemon
```

```
# api/Dockerfile.dev
FROM node:20-alpine

WORKDIR /app

COPY package*.json ./
RUN npm install

# Don't copy source - use volume mount
# COPY . .

EXPOSE 4000

CMD ["npm", "run", "dev"]
```

Development vs. Production Compose

docker-compose.yml (dev)

```
services:
  api:
    build:
      context: ./api
      dockerfile: Dockerfile.dev
    volumes:
      - ./api/src:/app/src
    environment:
      - NODE_ENV=development
    command: npm run dev
```

docker-compose.prod.yml

```
services:  
  api:  
    build:  
      context: ./api  
      dockerfile: Dockerfile  
    environment:  
      - NODE_ENV=production  
    restart: always  
    deploy:  
      replicas: 2
```

```
# Development  
docker compose up  
  
# Production  
docker compose -f docker-compose.prod.yml up
```

Debugging in Containers

```
services:  
  api:  
    build: ./api  
    ports:  
      - "4000:4000"  
      - "9229:9229" # Node.js debug port  
    command: node --inspect=0.0.0.0:9229 src/server.js
```

VS Code launch.json:

```
{  
  "type": "node",  
  "request": "attach",  
  "name": "Docker: Attach",  
  "port": 9229,  
  "address": "localhost",  
  "localRoot": "${workspaceFolder}/api/src",  
  "remoteRoot": "/app/src"  
}
```

Common Development Tasks

```
# Install new package
docker compose exec api npm install express-validator

# Run database migrations
docker compose exec api npm run migrate

# Access database shell
docker compose exec mongo mongosh myapp

# Run tests
docker compose exec api npm test

# View real-time logs
docker compose logs -f api

# Restart single service
docker compose restart api

# Rebuild and restart
docker compose up --build api
```

Live Demo

Building a MERN Stack with Docker Compose

Let's build this together:

1. Create project structure
2. Write Dockerfiles
3. Configure Docker Compose
4. Start the stack
5. Test hot reloading
6. Explore container networking

Project Structure

```
mern-app/
  docker-compose.yml
```

```
.env
frontend/
  Dockerfile
  Dockerfile.dev
  package.json
  src/
api/
  Dockerfile
  Dockerfile.dev
  package.json
  src/
.dockerignore
```

Step 1: API Dockerfile

```
# api/Dockerfile.dev
FROM node:20-alpine

WORKDIR /app

# Install dependencies
COPY package*.json .
RUN npm install

# Install nodemon globally for hot reload
RUN npm install -g nodemon

EXPOSE 4000

CMD ["nodemon", "src/server.js"]
```

Step 2: Docker Compose

```
# docker-compose.yml
services:
  api:
    build:
      context: ./api
      dockerfile: Dockerfile.dev
```

```

ports:
  - "4000:4000"
volumes:
  - ./api:/app
  - /app/node_modules
environment:
  - MONGODB_URI=mongodb://mongo:27017/mernapp
  - PORT=4000
depends_on:
  - mongo

mongo:
  image: mongo:7
  ports:
    - "27017:27017"
  volumes:
    - mongo-data:/data/db

mongo-express:
  image: mongo-express
  ports:
    - "8081:8081"
  environment:
    - ME_CONFIG_MONGODB_URL=mongodb://mongo:27017/
  depends_on:
    - mongo

volumes:
  mongo-data:

```

Step 3: Start and Verify

```

# Build and start
docker compose up --build

# In another terminal
curl http://localhost:4000/health

# Check MongoDB
docker compose exec mongo mongosh mernapp --eval "db.stats()"

```

```
# View logs
docker compose logs -f api

# Stop everything
docker compose down
```

Best Practices Summary

Docker Best Practices

1. Use official base images with specific versions
2. Minimize image size with Alpine variants and multi-stage builds
3. Don't run as root - use non-privileged users
4. Use .dockerignore to exclude unnecessary files
5. Leverage layer caching by ordering Dockerfile commands
6. Use health checks for production containers
7. Don't store secrets in images or environment variables
8. Tag images meaningfully - use semantic versioning

Compose Best Practices

1. Use named volumes for persistent data
2. Define networks explicitly for service isolation
3. Use depends_on with health checks for startup order
4. Separate dev and prod configurations
5. Use environment variables for configuration
6. Don't expose database ports in production
7. Set resource limits in production
8. Use restart policies appropriately

Looking Ahead

Lab 1: Next Week (January 20)

Infrastructure as Code with Docker Compose

You will:

- Set up a complete MERN development environment

- Write Dockerfiles for frontend and backend
- Configure multi-service orchestration
- Implement volume mounting for development
- Test the complete stack

Preparation for Lab 1

Install

- Docker Desktop
- VS Code with Docker extension
- Git

Verify

```
docker --version  
docker compose version
```

Review

- Today's slides
- [Docker Docs: Get Started](#)
- [Compose File Reference](#)

Summary

Today we covered:

1. Problems containers solve
2. Docker fundamentals - images, containers, volumes
3. Writing effective Dockerfiles
4. Multi-stage builds for optimization
5. Docker Compose for multi-service apps
6. Development workflows with containers
7. Container networking and data persistence

Questions?

Office Hours: Tuesday 9-11 AM, Pitt 2206

Email: kuruzj@rpi.edu

Appointments: bit.ly/jason-rpi

See you Tuesday for Lab 1!